
Project 5 - Can we predict whether a Hotel Booking will be canceled?

Ryan Chen 113200236

The Hotel Booking Dataset (Kaggle)

- Hotel Booking records for city and resort hotels
- Data collected from July 1, 2015 to August 31, 2017
- Dataset Size
 - 119390 rows
 - 36 columns



An Example Booking

```
hotel                Resort Hotel
is_canceled          0
lead_time             14
arrival_date_year     2015
arrival_date_month    July
arrival_date_week_number 27
arrival_date_day_of_month 1
stays_in_weekend_nights 0
stays_in_week_nights  2
adults               2
children             0.0
babies               0
meal                 BB
country              GBR
market_segment        Online TA
distribution_channel   TA/TO
is_repeated_guest     0
previous_cancellations 0
previous_bookings_not_canceled 0
reserved_room_type    A
assigned_room_type     A
booking_changes        0
deposit_type           No Deposit
```

```
agent                240.0
company              NaN
days_in_waiting_list 0
customer_type         Transient
adr                  98.0
required_car_parking_spaces 0
total_of_special_requests 1
reservation_status    Check-Out
reservation_status_date 2015-07-03
name                 Jasmine Fletcher
email                JFletcher43@xfinity.com
phone-number          190-271-6743
credit_card           *****9263
Name: 5, dtype: object
```

Exploratory Data Analysis

(Pandas, Numpy, Matplotlib)

Which country saw the most hotel bookings?

```
df.country.value_counts()
```

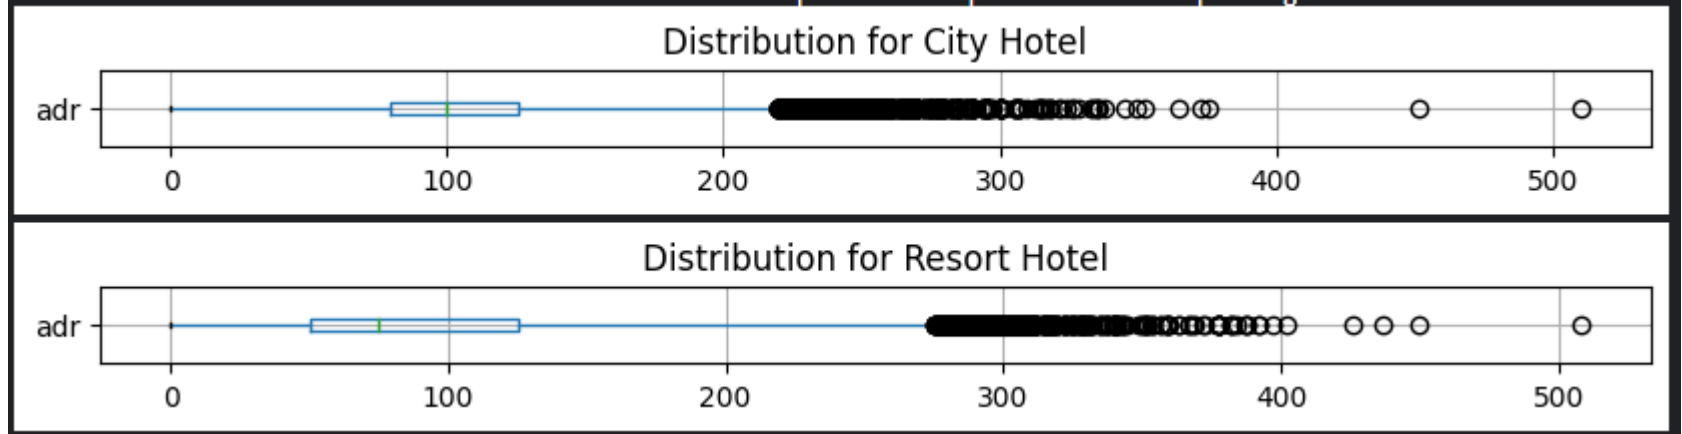
```
country
PRT    48590
GBR    12129
FRA    10415
ESP     8568
DEU     7287
...
DJI         1
BWA         1
HND         1
VGB         1
NAM         1
Name: count, Length: 177, dtype: int64
```



PRT = Portugal
accounts for 40.7 %
of all bookings in the
data

What is the distribution like for both hotels with respect to the price of a room per night?

column of interest = Average Daily Rate (adr)



Which months are most busy for both hotels?

```
df[df.hotel==h_type]  
.arrival_date_month  
.value_counts().  
head(3).index
```

City Hotel

- May
- July
- August

Resort Hotel

- April
 - July
 - August
-

Which months see the most expensive per night costs?

```
df[df.hotel==h_type]
.groupby('arrival_date_month')
['adr']
.mean().sort_values()
.tail(3)
```

```
The most expensive per night costs for City Hotel are:
arrival_date_month
August      114.68
June        119.07
May         121.64
Name: adr, dtype: float64
```

```
The most expensive per night costs for Resort Hotel are:
arrival_date_month
June        110.44
July        155.18
August      186.79
Name: adr, dtype: float64
```

Which months see the most cancellations for both hotel types?

```
df[df.hotel == h_type]
.groupby('arrival_date_month')
['is_canceled']
.sum().head()
```

```
For City Hotel , these months gave most cancellations
arrival_date_month
April          3465
August         3602
December       1740
February       1901
January        1482
Name: is_canceled, dtype: int64
```

```
For Resort Hotel , these months gave most cancellations
arrival_date_month
April          1059
August         1637
December        631
February        795
January         325
Name: is_canceled, dtype: int64
```

Examine distributions of bookings vs market segment

```
df.market_segment.value_counts()
```

```
Examine distributions of bookings vs market segment.  
market_segment  
Online TA      56477  
Offline TA/TO  24218  
Groups         19810  
Direct         12606  
Corporate       5295  
Complementary   743  
Aviation        237  
Undefined        2  
Name: count, dtype: int64
```

Which room type was most commonly booked? Most commonly cancelled?

```
reserved_room_type
A      85992
D      19201
E       6535
F       2897
G       2094
B       1118
C        932
H        601
P         12
L          6
Name: count, dtype: int64
```

```
reserved_room_type
A      33629
D       6102
E       1914
F        880
G        763
B        368
C        308
H        245
P         12
L          2
Name: is_canceled, dtype: int64
```

What percentage of the data recorded cancellations for each hotel?

`df[df.hotel==h_type]`

`.is_canceled.sum()*100`

`/ df[df.hotel==h_type].shape[0]`

Result

City Hotel bookings

- 41.73 % of bookings cancelled

Resort Hotel bookings


- 27.76 % of bookings cancelled

```
is_canceled
0      75166
1      44224
Name: count, dtype: int64
```

Machine Learning

What models to consider?

- Hotel Booking Cancellation is a classification task!
- A booking can be
 - canceled (1)
 - not canceled (0)



Data Cleaning (1)

- We have an overwhelming amount of features (36)
- Must eliminate some columns, initial dropping shown

```
['lead_time', 'arrival_date_year', 'arrival_date_week_number',  
 'arrival_date_day_of_month', 'previous_bookings_not_canceled',  
 'booking_changes', 'deposit_type', 'agent', 'company',  
 'required_car_parking_spaces', 'reservation_status',  
 'reservation_status_date', 'name', 'email', 'phone-number', 'credit_card']
```

Data Cleaning (2) - numerical cols

for col in ['adr', 'days_in_waiting_list']:

df[col] = np.log(df[col] + 0.000001)

| | | | | | |
|---------------------------|--------------|------------|----------|----------------|------------|
| is_canceled | Mean: 0.37 | STD: 0.48 | Min: 0.0 | Median: 0.0 | Max: 1.0 |
| stays_in_weekend_nights | Mean: 0.93 | STD: 1.0 | Min: 0.0 | Median: 1.0 | Max: 19.0 |
| stays_in_week_nights | Mean: 2.5 | STD: 1.91 | Min: 0.0 | Median: 2.0 | Max: 50.0 |
| adults | Mean: 1.86 | STD: 0.58 | Min: 0.0 | Median: 2.0 | Max: 55.0 |
| children | Mean: 0.1 | STD: 0.4 | Min: 0.0 | Median: 0.0 | Max: 10.0 |
| babies | Mean: 0.01 | STD: 0.1 | Min: 0.0 | Median: 0.0 | Max: 10.0 |
| is_repeated_guest | Mean: 0.03 | STD: 0.18 | Min: 0.0 | Median: 0.0 | Max: 1.0 |
| previous_cancellations | Mean: 0.09 | STD: 0.84 | Min: 0.0 | Median: 0.0 | Max: 26.0 |
| days_in_waiting_list | Mean: 2.32 | STD: 17.59 | Min: 0.0 | Median: 0.0 | Max: 391.0 |
| adr | Mean: 101.79 | STD: 48.15 | Min: 0.0 | Median: 94.575 | Max: 510.0 |
| total_of_special_requests | Mean: 0.57 | STD: 0.79 | Min: 0.0 | Median: 0.0 | Max: 5.0 |

Data Cleaning (3) - categorical vars

```
Col hotel has ['City Hotel' 'Resort Hotel'] unique values
Col arrival_date_month has ['April' 'August' 'December' 'February' 'January' 'July' 'June' 'March'
'May' 'November' 'October' 'September'] unique values
Col meal has ['BB' 'FB' 'HB' 'SC' 'Undefined'] unique values
Col country has ['ABW' 'AGO' 'AIA' 'ALB' 'AND' 'ARE' 'ARG' 'ARM' 'ASM' 'ATA' 'ATF' 'AUS'
'AUT' 'AZE' 'BDI' 'BEL' 'BEN' 'BFA' 'BGD' 'BGR' 'BHR' 'BHS' 'BIH' 'BLR'
'BOL' 'BRA' 'BRB' 'BWA' 'CAF' 'CHE' 'CHL' 'CHN' 'CIV' 'CMR' 'CN' 'COL'
'COM' 'CPV' 'CRI' 'CUB' 'CYM' 'CYP' 'CZE' 'DEU' 'DJI' 'DMA' 'DNK' 'DOM'
'DZA' 'ECU' 'EGY' 'ESP' 'EST' 'ETH' 'FIN' 'FJI' 'FRA' 'FRO' 'GAB' 'GBR'
'GEO' 'GGY' 'GHA' 'GIB' 'GLP' 'GNB' 'GRC' 'GTM' 'GUY' 'HKG' 'HND' 'HRV'
'HUN' 'IDN' 'IMN' 'IND' 'IRL' 'IRN' 'IRQ' 'ISL' 'ISR' 'ITA' 'JAM' 'JEY'
'JOR' 'JPN' 'KAZ' 'KEN' 'KHM' 'KIR' 'KNA' 'KOR' 'KWT' 'LAO' 'LBN' 'LBY'
'LCA' 'LIE' 'LKA' 'LTU' 'LUX' 'LVA' 'MAC' 'MAR' 'MCO' 'MDG' 'MDV' 'MEX'
'MKD' 'MLI' 'MLT' 'MMR' 'MNE' 'MOZ' 'MRT' 'MUS' 'MWI' 'MYS' 'MYT' 'NAM'
'NCL' 'NGA' 'NIC' 'NLD' 'NOR' 'NPL' 'NZL' 'OMN' 'PAK' 'PAN' 'PER' 'PHL'
'PLW' 'POL' 'PRI' 'PRT' 'PRY' 'PYF' 'QAT' 'ROU' 'RUS' 'RWA' 'SAU' 'SDN'
'SEN' 'SGP' 'SLE' 'SLV' 'SMR' 'SRB' 'STP' 'SUR' 'SVK' 'SVN' 'SWE' 'SYC'
'SYR' 'TGO' 'THA' 'TJK' 'TMP' 'TUN' 'TUR' 'TWN' 'TZA' 'UGA' 'UKR' 'UMI'
'URY' 'USA' 'UZB' 'VEN' 'VGB' 'VNM' 'ZAF' 'ZMB' 'ZWE'] unique values
Col market_segment has ['Aviation' 'Complementary' 'Corporate' 'Direct' 'Groups' 'Offline TA/TO'
'Online TA'] unique values
Col distribution_channel has ['Corporate' 'Direct' 'GDS' 'TA/TO' 'Undefined'] unique values
Col reserved_room_type has ['A' 'B' 'C' 'D' 'E' 'F' 'G' 'H' 'L' 'P'] unique values
Col assigned_room_type has ['A' 'B' 'C' 'D' 'E' 'F' 'G' 'H' 'I' 'K' 'L' 'P'] unique values
Col customer_type has ['Contract' 'Group' 'Transient' 'Transient-Party'] unique values
```

How to deal with categorical data?

- Many ML models can only accept numerical data
 - Solution = use encoding techniques
 - One Hot Encoding
 - Ordinal Encoding
-

```
df_encoded = pd.get_dummies(df, columns=['hotel', 'meal', 'market_segment',  
                                         'distribution_channel', 'reserved_room_type',  
                                         'assigned_room_type', 'customer_type'],  
                             drop_first=True) # drop_first=True to avoid dummy variable trap
```

```
# ordinal encode the arrival_date_month column  
from sklearn.preprocessing import OrdinalEncoder  
df = df_encoded  
month_order = ['January', 'February', 'March', 'April', 'May', 'June',  
               'July', 'August', 'September', 'October', 'November', 'December']  
ordinal_encoder = OrdinalEncoder(categories=[month_order])  
df['arrival_date_month_encoded'] = ordinal_encoder.fit_transform(df[['arrival_date_month']])  
df = df.drop(columns=['arrival_date_month'])  
print(df.iloc[5])
```

The resulting columns shown

```
is_canceled          0
stays_in_weekend_nights  0
stays_in_week_nights  2
adults               2
children             0.0
babies              0
is_repeated_guest    0
previous_cancellations 0
days_in_waiting_list -13.815511
adr                  4.584967
total_of_special_requests 1
hotel_Resort Hotel    True
meal_FB              False
meal_HB              False
meal_SC              False
meal_Undefined        False
market_segment_Complementary False
market_segment_Corporate False
market_segment_Direct False
market_segment_Groups False
market_segment_Offline TA/TO False
market_segment_Online TA True
distribution_channel_Direct False
distribution_channel_GDS False
distribution_channel_TA/TO True
distribution_channel_Undefined False
```

```
reserved_room_type_B False
reserved_room_type_C False
reserved_room_type_D False
reserved_room_type_E False
reserved_room_type_F False
reserved_room_type_G False
reserved_room_type_H False
reserved_room_type_L False
reserved_room_type_P False
assigned_room_type_B False
assigned_room_type_C False
assigned_room_type_D False
assigned_room_type_E False
assigned_room_type_F False
assigned_room_type_G False
assigned_room_type_H False
assigned_room_type_I False
assigned_room_type_K False
assigned_room_type_L False
assigned_room_type_P False
customer_type_Group False
customer_type_Transient True
customer_type_Transient-Party False
arrival_date_month_encoded 6.0
Name: 5, dtype: object
```

Which classifiers should we choose?

- 1) Logistic Regression
- 2) Decision Tree
- 3) Gradient Boosting Classifier

* sklearn is our library of choice here!

But first Train Test Split

```
# split the data into train and test sets
from sklearn.model_selection import train_test_split
X = df.drop(columns=['is_canceled'])
y = df['is_canceled']
assert 'is_canceled' not in X.columns
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.33, random_state=42)
```


Model Training

```
# decision tree model
from sklearn import tree
decision_tree = tree.DecisionTreeClassifier()
decision_tree.fit(X_train,y_train)
y_pred = decision_tree.predict(X_test)
```

```
# Logistic Regression Model
from sklearn.linear_model import LogisticRegression
lr_model = LogisticRegression(max_iter=1000)
lr_model.fit(X_train, y_train)
y_pred = lr_model.predict(X_test)
print(y_pred)
```

```
# gradient boosting classifier
from sklearn.ensemble import GradientBoostingClassifier
grad_boost_model = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0,
                                              max_depth=1, random_state=0).fit(X_train, y_train)
y_pred = grad_boost_model.predict(X_test)
```

Classification Performance



How to compute these measures?

```
print("The decision tree model results")
acc = metrics.accuracy_score(y_test, y_pred)
print("Accuracy is ",round(acc,2))
precision = metrics.precision_score(y_test, y_pred)
print("Precision is ",round(precision,2))
recall = metrics.recall_score(y_test, y_pred)
print("Recall is ",round(recall,2))
f1_score = metrics.f1_score(y_test, y_pred)
print('F1 score is ',round(f1_score,2))
```

Overall results + Winner!!

| Model | Accuracy | Precision | Recall | F1 Score |
|---------------|----------|-----------|--------|----------|
| Log Reg | 0.77 | 0.73 | 0.6 | 0.66 |
| Decision Tree | 0.79 | 0.72 | 0.72 | 0.72 |
| Grad Boost | 0.78 | 0.75 | 0.59 | 0.67 |

Additional Considerations....

- Additional Feature Reduction
 - More models to consider
 - Consideration of hyperparameter tuning for each model
-

